

THE WORLD OF SELECT()

So just why am I so hyped on select()?

One traditional way to write network servers is to have the main server block on `accept()`, waiting for a connection. Once a connection comes in, the server `fork()`s, the child process handles the connection and the main server is able to service new incoming requests.

With `select()`, instead of having a process for each request, there is usually only one process that "multi-plexes" all requests, servicing each request as much as it can.

So one main advantage of using `select()` is that your server will only require a single process to handle all requests. Thus, your server will not need shared memory or synchronization primitives for different 'tasks' to communicate.

One major disadvantage of using `select()`, is that your server cannot act like there's only one client, like with a `fork()`'ing solution. For example, with a `fork()`'ing solution, after the server `fork()`s, the child process works with the client as if there was only one client in the universe -- the child does not have to worry about new incoming connections or the existence of other sockets. With `select()`, the programming isn't as transparent.

Okay, so how do you use select()?

`select()` works by blocking until something happens on a file descriptor (aka a socket). What's 'something'? Data coming in or being able to write to a file descriptor -- you tell `select()` what you want to be woken up by. How do you tell it? You fill up a `fd_set` structure with some macros.

Most `select()`-based servers look pretty much the same:

- Fill up a `fd_set` structure with the file descriptors you want to know when data comes in on.
- Fill up a `fd_set` structure with the file descriptors you want to know when you can write on.
- Call `select()` and block until something happens.
- Once `select()` returns, check to see if any of your file descriptors was the reason you woke up. If so, 'service' that file descriptor in whatever particular way your server needs to (i.e. read in a request for a Web page).
- Repeat this process forever.

Quit with the pseudo-code, show me some real code!

Okay, let's take a look at a [sample server \(original\)](#) included with Vic Metcalfe's [Socket Programming FAQ](#) (my comments are in **red**):

```
/*
```

```
PLEASE READ THE FILE NB-APOLOGY!!!! There are some things you should
know about this source before you read it. Thanks.
```

```
Quang Ngo alerted me to a bug where the variable listnum in deal_with_data()
wasn't being passed in by parameter, thus it was always garbage. I have
quick-fixed this in the code below. - Spencer (October 12, 1999)
```

Non blocking server demo
 By Vic Metcalfe (vic@acm.org)
 For the [unix-socket-faq](#)

```

*/

#include "sockhelp.h"
#include <ctype.h>
#include <sys/time.h>
#include <fcntl.h>

int sock;          /* The socket file descriptor for our "listening"
                    socket */
int connectlist[5]; /* Array of connected sockets so we know who
                    we are talking to */
fd_set socks;      /* Socket file descriptors we want to wake
                    up for, using select() */
int highsock;      /* Highest #'d file descriptor, needed for select() */

void setnonblocking(sock)
int sock;
{
    int opts;

    opts = fcntl(sock,F_GETFL);
    if (opts < 0) {
        perror("fcntl(F_GETFL)");
        exit(EXIT_FAILURE);
    }
    opts = (opts | O_NONBLOCK);
    if (fcntl(sock,F_SETFL,opts) < 0) {
        perror("fcntl(F_SETFL)");
        exit(EXIT_FAILURE);
    }
    return;
}

void build_select_list() {
    int listnum;          /* Current item in connectlist for for loops */

    /* First put together fd_set for select(), which will
       consist of the sock variable in case a new connection
       is coming in, plus all the sockets we have already
       accepted. */

    /* FD_ZERO() clears out the fd_set called socks, so that
       it doesn't contain any file descriptors. */

    FD_ZERO(&socks);

    /* FD_SET() adds the file descriptor "sock" to the fd_set,
       so that select() will return if a connection comes in
       on that socket (which means you have to do accept(), etc. */

    FD_SET(sock,&socks);

    /* Loops through all the possible connections and adds
       those sockets to the fd_set */

    for (listnum = 0; listnum < 5; listnum++) {
        if (connectlist[listnum] != 0) {
            FD_SET(connectlist[listnum],&socks);
            if (connectlist[listnum] > highsock)
                highsock = connectlist[listnum];
        }
    }
}

void handle_new_connection() {

```

```

int listnum;          /* Current item in connectlist for for loops */
int connection; /* Socket file descriptor for incoming connections */

/* We have a new connection coming in! We'll
try to find a spot for it in connectlist. */
connection = accept(sock, NULL, NULL);
if (connection < 0) {
    perror("accept");
    exit(EXIT_FAILURE);
}
setnonblocking(connection);
for (listnum = 0; (listnum < 5) && (connection != -1); listnum++)
    if (connectlist[listnum] == 0) {
        printf("\nConnection accepted:  FD=%d; Slot=%d\n",
            connection, listnum);
        connectlist[listnum] = connection;
        connection = -1;
    }
if (connection != -1) {
    /* No room left in the queue! */
    printf("\nNo room left for new client.\n");
    sock_puts(connection, "Sorry, this server is too busy.  "
        "Try again later!\n\n");
    close(connection);
}
}

void deal_with_data(
    int listnum          /* Current item in connectlist for for loops */
) {
    char buffer[80];      /* Buffer for socket reads */
    char *cur_char;       /* Used in processing buffer */

    if (sock_gets(connectlist[listnum], buffer, 80) < 0) {
        /* Connection closed, close this end
        and free up entry in connectlist */
        printf("\nConnection lost: FD=%d; Slot=%d\n",
            connectlist[listnum], listnum);
        close(connectlist[listnum]);
        connectlist[listnum] = 0;
    } else {
        /* We got some data, so upper case it
        and send it back. */
        printf("\nReceived: %s; ", buffer);
        cur_char = buffer;
        while (cur_char[0] != 0) {
            cur_char[0] = toupper(cur_char[0]);
            cur_char++;
        }
        sock_puts(connectlist[listnum], buffer);
        sock_puts(connectlist[listnum], "\n");
        printf("responded: %s\n", buffer);
    }
}

void read_socks() {
    int listnum;          /* Current item in connectlist for for loops */

    /* OK, now socks will be set with whatever socket(s)
    are ready for reading. Lets first check our
    "listening" socket, and then check the sockets
    in connectlist. */

    /* If a client is trying to connect() to our listening
    socket, select() will consider that as the socket
    being 'readable'. Thus, if the listening socket is
    part of the fd_set, we need to accept a new connection. */

    if (FD_ISSET(sock, &socks))

```

```

        handle_new_connection();
    /* Now check connectlist for available data */

    /* Run through our sockets and check to see if anything
       happened with them, if so 'service' them. */

    for (listnum = 0; listnum < 5; listnum++) {
        if (FD_ISSET(connectlist[listnum],&socks))
            deal_with_data(listnum);
    } /* for (all entries in queue) */
}

int main (argc, argv)
int argc;
char *argv[];
{
    char *ascpport; /* ASCII version of the server port */
    int port; /* The port number after conversion from ascpport */
    struct sockaddr_in server_address; /* bind info structure */
    int reuse_addr = 1; /* Used so we can re-bind to our port
                        while a previous connection is still
                        in TIME_WAIT state. */
    struct timeval timeout; /* Timeout for select */
    int readsocks; /* Number of sockets ready for reading */

    /* Make sure we got a port number as a parameter */
    if (argc < 2) {
        printf("Usage: %s PORT\r\n",argv[0]);
        exit(EXIT_FAILURE);
    }

    /* Obtain a file descriptor for our "listening" socket */
    sock = socket(AF_INET, SOCK_STREAM, 0);
    if (sock < 0) {
        perror("socket");
        exit(EXIT_FAILURE);
    }
    /* So that we can re-bind to it without TIME_WAIT problems */
    setsockopt(sock, SOL_SOCKET, SO_REUSEADDR, &reuse_addr,
               sizeof(reuse_addr));

    /* Set socket to non-blocking with our setnonblocking routine */
    setnonblocking(sock);

    /* Get the address information, and bind it to the socket */
    ascpport = argv[1]; /* Read what the user gave us */
    port = atoi(port); /* Use function from sockhelp to
                        convert to an int */
    memset((char *) &server_address, 0, sizeof(server_address));
    server_address.sin_family = AF_INET;
    server_address.sin_addr.s_addr = htonl(INADDR_ANY);
    server_address.sin_port = port;
    if (bind(sock, (struct sockaddr *) &server_address,
              sizeof(server_address)) < 0 ) {
        perror("bind");
        close(sock);
        exit(EXIT_FAILURE);
    }

    /* Set up queue for incoming connections. */
    listen(sock,5);

    /* Since we start with only one socket, the listening socket,
       it is the highest socket so far. */
    highsock = sock;
    memset((char *) &connectlist, 0, sizeof(connectlist));

    while (1) { /* Main server loop - forever */
        build_select_list();

```

```

timeout.tv_sec = 1;
timeout.tv_usec = 0;

/* The first argument to select is the highest file
   descriptor value plus 1. In most cases, you can
   just pass FD_SETSIZE and you'll be fine. */

/* The second argument to select() is the address of
   the fd_set that contains sockets we're waiting
   to be readable (including the listening socket). */

/* The third parameter is an fd_set that you want to
   know if you can write on -- this example doesn't
   use it, so it passes 0, or NULL. The fourth parameter
   is sockets you're waiting for out-of-band data for,
   which usually, you're not. */

/* The last parameter to select() is a time-out of how
   long select() should block. If you want to wait forever
   until something happens on a socket, you'll probably
   want to pass NULL. */

readsocks = select(highsock+1, &socks, (fd_set *) 0,
                  (fd_set *) 0, &timeout);

/* select() returns the number of sockets that had
   things going on with them -- i.e. they're readable. */

/* Once select() returns, the original fd_set has been
   modified so it now reflects the state of why select()
   woke up. i.e. If file descriptor 4 was originally in
   the fd_set, and then it became readable, the fd_set
   contains file descriptor 4 in it. */

if (readsocks < 0) {
    perror("select");
    exit(EXIT_FAILURE);
}
if (readsocks == 0) {
    /* Nothing ready to read, just show that
       we're alive */
    printf(".");
    fflush(stdout);
} else
    read_socks();
} /* while(1) */
} /* main */

```

Okay, so maybe that wasn't the best example...

Got some suggests? Corrections? Please let me know.

In the mean time, here's some references to other sources of information about select():

- [Socket FAQ Question about select\(\).](#)
Straight from the excellent socket FAQ.
- [Unix Programming FAQ Question about select\(\).](#)
Straight from the *other* excellent FAQ.
- [Unix Socket Programming FAQ Examples](#)
The above nbsvr.c sample code and more socket stuff.
- [thttpd - tiny/turbo/throttling HTTP server](#)
Nifty little single-threaded, non-fork()'ing, select() based Web server.
- [BOA](#)

Another single-threaded, select() based Web server.



[Back to Spencer's Socket Site](#)

[Copyright](#) © 1995-2022 [Spencer Low](#). All rights reserved. [Privacy Policy](#)